



02456 – Week 4


Convolutional

Neural Networks

Jes Frellsen

Technical University of Denmark

22 September 2025



Menu of the day

Basic supervised
learning with NNs

- Week 1: Neural nets
- Week 2: Learning
- Week 3: Tricks of The Trade

Specialised
architectures


- **Week 4: CNNs**
- Week 5: RNNs
- Week 6: Transformers
- Week 7: Unsupervised
- Week 8: Mini-project

Lecture


- Motivation for CNNs
- Details of CNNs (ch 10)
- Applications of CNNs (ch 10)

Exercises

- **Notebooks 4.1-4.3:**
 - CNNs in **PyTorch**!
- **Problems:** 10.1, 10.6, 10.16



Recap: Neural networks and learning



We can write the **output of layer ℓ** as

$$f^{(\ell)}(\mathbf{h}) = \sigma(W^{(\ell)}\mathbf{h} + \mathbf{b}^{(\ell)}),$$

where σ is a non-linear **activation function**.

The **joint function** is then



$$f_{\phi}(\mathbf{x}) = \mathbf{y} = f^{(4)}\left(f^{(3)}\left(f^{(2)}\left(f^{(1)}(\mathbf{x})\right)\right)\right)$$

A **loss function** $L(\phi)$ captures the **mismatch** of $f_{\phi}(\mathbf{x}_i)$ and y_i : $\hat{\phi} = \underset{\phi}{\operatorname{argmin}} L(\phi)$


Use the **negative log-likelihood**

$$L(\phi) = -\sum_i \log p(y_i | f_{\phi}(\mathbf{x}_i))$$

- Use **gradient descent** to find $\hat{\phi}$
 - Backpropagation** gets gradients $\nabla_{\phi} L(\phi)$
 - Initialisation** is important
- Split **data set** in three
 - Val**: Select hyper parameters
 - Test**: estimate generalisation error




Motivation: Image classification



ImageNet

- 1,000 classes
 - Including many types of dogs!
- ~1.2M training images
- Performance measure: is the correct class within top, e.g., 5% (due to ambiguity)

Convolutional neural networks



$$\begin{bmatrix} 10 & 0 & -10 \\ 0 & 0 & 0 \\ -10 & 0 & 10 \end{bmatrix}$$




High-level idea:

- **Learn to extract features** from image
- **Same weights** applied across positions
 - Weight sharing
 - Translation equivariant
- **Sub-sampling (pooling)**
 - Reduce spatial dim.
 - Increase num. features
- **MLP at the end**

Feature engineering vs engineered models

ImageNet Classification with Deep Convolutional Neural Networks


ILSVRC top-5 error on ImageNet




Alex Krizhevsky
University of Toronto
kriz@cs.utoronto.ca

Ilya Sutskever
University of Toronto
ilya@cs.utoronto.ca


Geoffrey E. Hinton
University of Toronto
hinton@cs.utoronto.ca



AlexNet: Learned filters in first layer



Emergent higher-level abstractions



Activations of the 151st channel of **layer 5** in deep CNN:

- Sensitive to **faces**

Remember **learned!**


- We **designs** an **architecture** that make the network **learn increasingly abstract filters**

Convolutional Networks


Issues with Fully Connected Layers for Images

- **First weight matrix** has size $|\text{input}| \times |\mathbf{h}^{(1)}|$
 - For a 1000×1000 image: $10^6 \times |\mathbf{h}^{(1)}|$ parameters
- **No notion of locality** (“nearby” pixels are not treated specially)
 - Permuting pixels + permuting weights leaves output unchanged
- Interpretation of an image **not stable under translations**
 - We want layers that are **equivariant to translations**
- **Convolutional layers address all these issues**
 - Popularized by **Yann LeCun** et al. (1989)
 - Two ideas: **Local connectivity** and **parameter sharing**

Local connectivity




25 parameters




13 parameters

Parameter sharing







13 parameters



3 parameters

If we learn to extract a feature on place in the image, we can extract it in all places!

Convolution *




The **convolution*** on a 2D image $\mathbf{X}$ with a 2D kernel/filter $\mathbf{W}$

$$z_{ij} = (\mathbf{X} * \mathbf{W}) = \sum_m \sum_n X_{i+m, j+n} W_{m,n}$$

*) Strickly speaking a **cross-correlation**

Convolutional Neural Networks



We use **convs** to construct a **CNN**

- Pooling to down-sample
- Fully connected layer at the end

Central: Parameter **sharing through filters**

- By sliding a filter give **feature map**
- **Multiple filters** give **multiple channels**
- Feature maps gets **increasingly abstract** away for the input

What does a filter do?

filter

0	0	1
0	0	1
0	0	1

We can view the operation as an inner product

image

1	0	0	0	0	0
1	1	0	1	1	0
0	0	0	0	1	1
0	0	0	1	0	0
1	0	0	1	0	0
1	1	0	1	0	0

result

0	1	2	1
0	2	2	1
0	2	1	1
0	3	0	0

Pooling?

image

1	0	0	0	0
3	2	0	1	1
0	0	0	0	2
0	0	0	1	0
1	0	0	3	0

avg pool

$\frac{6}{9}$	$\frac{4}{9}$
$\frac{1}{9}$	$\frac{6}{9}$

Invariance and equivalence


A function f is **invariant** to a transformation t if

$$f(t(\mathbf{x})) = f(\mathbf{x})$$


A function f is **equivariant** to a transformation t if

$$f(t(\mathbf{x})) = t(f(\mathbf{x}))$$



In images we want **translational equivalence**



Conv layers are translational equivariant



Example: CNN on MNIST



- MNIST data: 28×28 pixels and 1 channel
- First hidden layer: 28×28 pixels and 9 channels
- Use 5×5 filters

Notation for convolutional layer

- Replace **vector** signals with **tensors** indexed by (x, y, c)
 - x, y are spatial coordinates and c is the channel
- Convolutions weights is also tensor $\mathbf{W} \in \mathbb{R}^{w \times h \times c_{\text{in}} \times c_{\text{out}}}$

$$h_{x,y,c_1}^{(j)} = \sigma \left(\sum_{c_0, m, n} h_{x+m, y+n, c_0}^{(j-1)} W_{m, n, c_0, c_1} + b_{c_1} \right)$$




or

$$\mathbf{h}_{:, :, c_1}^{(j)} = \sigma \left(\sum_{c_0} \mathbf{h}_{:, :, c_0}^{(j-1)} * W_{:, :, c_0, c_1} + b_{c_1} \right)$$

Padding: Input can be padded with zeroes to keep the same spatial size of the input and output.

Pooling



- Decrease resolution (pooling)
 - Increase channels going up (convolutions).
- Often down-sampling by hard-coded pooling layers.
- Here, use $\max(\cdot)$ of 2×2 activations





$\mathbf{h}^{(1)}$ before and after pooling

Stack more layers


Note how each unit looks at all 9 channels of the previous layer



Whole network





8



- At the top, the resolution drops to 1×1 pixel
- Train by back-propagation as usual

Good learned representations?



- In pixel space, 9-4 is closer.
- In feature space, 4-4 is closer.

Convolutional vs. non-convolutional

CNN example:

- Sizes of signals/features **h**:

$$28 \cdot 28 = 784$$

$$28 \cdot 28 \cdot 9 = 7,056$$

$$14 \cdot 14 \cdot 16 = 3,136$$

- Number of weights (ignoring biases):

$$5 \cdot 5 \cdot 9 + 5 \cdot 5 \cdot 9 \cdot 16 + 7 \cdot 7 \cdot 16 \cdot 10 = 225 + 3600 + 7840 = 11,665$$

Compare to a MNIST FFNN example:

- Signals/features:


$$784, 225, 144$$

- Weights:

$$28 \cdot 28 \cdot 225 + 225 \cdot 144 + 144 \cdot 10 = 176400 + 32400 + 1440 = 210,240$$

Convolutional network has more signals/features but less parameter

Scaling up



- **Trend in 2015 towards:**


- Small filters (3×3)
- Smaller pooling (2×2)
- More layers (10 – 40)

- **Trends 2015+:**


- Even smaller filters, sometimes 1×1 ,
- Even more layers
- Average pooling in transition from convolutional to fully connected layers

Deep residual nets

- **Observation:** it is difficult to fit very deep nets





- **Solution:** skip layers $\mathbf{h}^{(\ell)} = \mathbf{h}^{(\ell-1)} + f_{\phi}(\mathbf{h}^{(\ell)})$




- Initially $f_{\phi}(\mathbf{h}^{(\ell)}) \approx 0$ meaning net $\approx$ linear

Deep residual nets: ImageNet results



method	top-1 err.	top-5 err.
VGG [41] (ILSVRC'14)	-	8.43 [†]
GoogLeNet [44] (ILSVRC'14)	-	7.89
VGG [41] (v5)	24.4	7.1
PRReLU-net [13]	21.59	5.71
BN-inception [16]	21.99	5.81
ResNet-34 B	21.84	5.71
ResNet-34 C	21.53	5.60
ResNet-50	20.74	5.25
ResNet-101	19.87	4.60
ResNet-152	19.38	4.49

Application: Semantic segmentation



- Use a CNN to classify each pixel
- Lots of shared computation


Ground Truth




Image



FCN-8s



U-net: Utilising residual connections for segmentation



Today's exercises!

- Three Python notebooks ([Week4 Convolutional](#))
 - Implement CNNs in **PyTorch** MNIST and CIFAR-10
 - Transfer learning from CIFAR-10 to the Caltech101 datasets
- Three **problems** from to book on covered material

Thank you!